
recoverability-aware planning (3.5× the tokens of one heal, 21 cheap dev evaluations)
1. Prune
remove the k blocks with the lowest Block Influence
2. Diagnose
retention $r_c(M_p)$ on six capabilities (dev split)
3. Pilot heals
6 pools × 3 budgets + 3 scopes (dev)
4. Fit
$g_c(x) = a_c(1 - e^{-x/\tau_c})$
transfer $T \geq 0$
5. Allocate
concave program: mixture $\mathbf{t}^\star$ and scope $\sigma^\star$
6. Heal once
B tokens with $(\mathbf{t}^\star, \sigma^\star)$; evaluate on test
teacher = unpruned model (self-distillation targets)
safeguards: τ lower bound, trust region, capped objective
baselines replace this step: uniform / damage-proportional